

Exercise 2b - MLP

We created a script that implements a Multi-Layer Perceptron with a variable-size hidden layer, and ran it on the 60000 samples of the provided MNIST Dataset. We settled, after testing, on a learning rate of 0.1, 60 neurons in the hidden layer, and a 100 training epochs, to obtain an accuracy of 99.3% and average loss of 0.026487 on the training set, and an accuracy of 80.5% and average loss of 0.688914 on the testing set. The plots available under this paragraph show, for both training and testing sets, the accuracies and losses at each epoch.

